


```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import nltk
6 from nltk.stem.porter import PorterStemmer
7 nltk.download('stopwords')
8 from nltk.corpus import stopwords
9 STOPWORDS = set(stopwords.words('english'))
10
11 from sklearn.model_selection import train_test_split
12 from sklearn.preprocessing import MinMaxScaler
13 from sklearn.feature_extraction.text import CountVectorizer
14 from sklearn.model_selection import cross_val_score
15 from sklearn.ensemble import RandomForestClassifier
16 from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
17 from sklearn.model_selection import GridSearchCV
18 from sklearn.model_selection import StratifiedKFold
19 from sklearn.metrics import accuracy_score
20 from wordcloud import WordCloud
21 from sklearn.tree import DecisionTreeClassifier
22 from xgboost import XGBClassifier
23 import pickle
24 import re

```

 [nltk\_data] Downloading package stopwords to /root/nltk\_data...  
 [nltk\_data] Unzipping corpora/stopwords.zip.


```
1 %pip install wordcloud
```

 Requirement already satisfied: wordcloud in /usr/local/lib/python3.10/dist-packages (1.9.3)  
 Requirement already satisfied: numpy>=1.6.1 in /usr/local/lib/python3.10/dist-packages (from wordcloud) (1.26.4)  
 Requirement already satisfied: pillow in /usr/local/lib/python3.10/dist-packages (from wordcloud) (10.4.0)  
 Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (from wordcloud) (3.7.1)  
 Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (1.3.0)  
 Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (0.12.1)  
 Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (4.54.1)  
 Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (1.4.7)  
 Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (24.1)  
 Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (3.2.0)  
 Requirement already satisfied: python-dateutil>=2.7 in /usr/local/lib/python3.10/dist-packages (from matplotlib->wordcloud) (2.8.2)  
 Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.10/dist-packages (from python-dateutil>=2.7->matplotlib->wordcloud) (1.16.0)


```

1 #Load the data
2
3 data = pd.read_csv(r"/content/2111.txt", delimiter = '\t', quoting = 3)
4
5 print(f"Dataset shape : {data.shape}")
6

```

 Dataset shape : (2769, 5)

```
1 data.head()
```




	feedback	variation	date	rating	verified_reviews
0	0	Mamaearth-Onion-Growth-Control-Redensyl	2019-09-06 00:00:00	1	I bought this hair oil after viewing so many g...
1	1	Mamaearth-Onion-Growth-Control-Redensyl	2019-08-14 00:00:00	5	"Used This Mama Earth Newly Launched Onion Oil...
2	0	Mamaearth-Onion-Growth-Control-Redensyl	2019-10-19 00:00:00	1	So bad product...My hair falling increase too ...
3	0	Mamaearth-Onion-Growth-Control-Redensyl	2019-09-16 00:00:00	1	Product just smells similar to navarathna hair...
4	1	Mamaearth-Onion-Growth-Control-Redensyl	2019-08-18 00:00:00	5	I have been trvina different onion oil for mv ...

Next steps: [Generate code with data](#) [View recommended plots](#) [New interactive sheet](#)

```

1 #Column names
2
3 print(f"Feature names : {data.columns.values}")


```

 Feature names : ['feedback' 'variation' 'date' 'rating' 'verified\_reviews']

```


1 #Check for null values
2
3 data.isnull().sum()

```




	0
feedback	0
variation	0
date	0
rating	0
verified_reviews	6

```
1
2 data.isnull().sum()
```



	0
feedback	0
variation	0
date	0
rating	0
verified_reviews	6

```
1 #Getting the record where 'verified_reviews' is null
2
3 data[data['verified_reviews'].isna() == True]
```



	feedback	variation	date	rating	verified_reviews
2238	1	Tata-Tea-Gold-500g	2018-03-03 00:00:00	4	NaN
2248	1	Tata-Tea-Gold-500g	2018-03-03 00:00:00	4	NaN
2621	1	Mysore-Sandal-Bathing-Soap-125g	2020-05-22 00:00:00	4	NaN
2624	1	Mysore-Sandal-Bathing-Soap-125g	2020-09-24 00:00:00	5	NaN
2631	1	Mysore-Sandal-Bathing-Soap-125g	2020-05-22 00:00:00	4	NaN
2634	1	Mysore-Sandal-Bathing-Soap-125g	2020-09-24 00:00:00	5	NaN


```
1 #We will drop the null record
2
3 data.dropna(inplace=True)
```

```
1 print(f"Dataset shape after dropping null values : {data.shape}")
```

 Dataset shape after dropping null values : (2763, 5)

```
1 data['length'] = data['verified_reviews'].apply(len)
```

```
1 data.head()
```



	feedback	variation	date	rating	verified_reviews	length
0	0	Mamaearth-Onion-Growth-Control-Redensyl	2019-09-06 00:00:00	1	I bought this hair oil after viewing so many g...	477
1	1	Mamaearth-Onion-Growth-Control-Redensyl	2019-08-14 00:00:00	5	"Used This Mama Earth Newly Launched Onion Oil...	497
2	0	Mamaearth-Onion-Growth-Control-Redensyl	2019-10-19 00:00:00	1	So bad product...My hair falling increase too ...	149

Next steps:

[Generate code with data](#)[View recommended plots](#)[New interactive sheet](#)

```
1 #Randomly checking for 10th record
2
3 print(f"'verified_reviews' column value: {data.iloc[10]['verified_reviews']}") #Original value
4 print(f"Length of review : {len(data.iloc[10]['verified_reviews'])}") #Length of review using len()
5 print(f"'length' column value : {data.iloc[10]['length']}") #Value of the column 'length'
```

```

↵ 'verified_reviews' column value: I bought this hair oil after viewing so many good comments. But this product is not good enough.Fir
Length of review : 477
'length' column value : 477

```

```
1 data.dtypes
```

```

↵
      0
feedback    int64
variation    object
date         object
rating       int64
verified_reviews object
length       int64

```

analysing rating column

```
1 len(data)
```

```
↵ 2763
```

```
1 print(f"Rating value count: \n{data['rating'].value_counts()}")
```

```

↵ Rating value count:
rating
5    1433
1     544
4     460
3     198
2     128
Name: count, dtype: int64


```

Let's plot the above values in a bar graph

```

1 #Bar plot to visualize the total counts of each rating
2
3 data['rating'].value_counts().plot.bar(color = 'red')
4 plt.title('Rating distribution count')
5 plt.xlabel('Ratings')
6 plt.ylabel('Count')
7 plt.show()

```



```

1
2 #Finding the percentage distribution of each rating - we'll divide the number of records for each rating by total number of records
3
4 print(f"Rating value count - percentage distribution: \n{round(data['rating'].value_counts()/data.shape[0]*100,2)}")

```

```


Rating value count - percentage distribution:
rating
5    51.86
1    19.69
4    16.65
3     7.17
2     4.63
Name: count, dtype: float64

```

```

1
2 fig = plt.figure(figsize=(7,7))
3
4 colors = ('pink', 'green', 'blue','orange','yellow')
5
6 wp = {'linewidth':1, "edgecolor":'black'}
7
8 tags = data['rating'].value_counts()/data.shape[0]
9
10 explode=(0.1,0.1,0.1,0.1,0.1)
11
12 tags.plot(kind='pie', autopct="%1.1f%%", shadow=True, colors=colors, startangle=90, wedgeprops=wp, explode=explode, label='Percentage
13
14 from io import BytesIO
15
16 graph = BytesIO()
17
18 fig.savefig(graph, format="png")

```



Analyzing 'feedback' column

```

1 #Distinct values of 'feedback' and its count
2
3 print(f"Feedback value count: \n{data['feedback'].value_counts()}")

```

```

Feedback value count:
feedback
1    2091
0     672
Name: count, dtype: int64

```

```

1 #Extracting the 'verified_reviews' value for one record with feedback = 0
2
3 review_0 = data[data['feedback'] == 0].iloc[1]['verified_reviews']
4 print(review_0)

```

```

So bad product...My hair falling increase too much...I order shampoo mask and oil.. nothing stop hairfallAfter 3 to 4 wash my hair fi

```

```

1 #Extracting the 'verified_reviews' value for one record with feedback = 1
2
3 review_1 = data[data['feedback'] == 1].iloc[1]['verified_reviews']
4 print(review_1)

```


I have been trying different onion oil for my hair as my hair is not very healthy. This product has literally changed the texture of

From the above 2 examples we can see that feedback 0 is negative review and 1 is positive review

```

1 #Bar graph to visualize the total counts of each feedback
2
3 data['feedback'].value_counts().plot.bar(color = 'blue')
4 plt.title('Feedback distribution count')
5 plt.xlabel('Feedback')
6 plt.ylabel('Count')
7 plt.show()

```



```

1 #Finding the percentage distribution of each feedback - we'll divide the number of records for each feedback by total number of records
2
3 print(f"Feedback value count - percentage distribution: \n{round(data['feedback'].value_counts()/data.shape[0]*100,2)}")

```

Feedback value count - percentage distribution:

feedback	count	percentage
1	2050	75.68
0	650	24.32


Name: count, dtype: float64

```

1 fig = plt.figure(figsize=(7,7))
2
3 colors = ('pink', 'skyblue')
4
5 wp = {'linewidth':1, "edgecolor":'black'}
6
7 tags = data['feedback'].value_counts()/data.shape[0]
8
9 explode=(0.1,0.1)
10
11 tags.plot(kind='pie', autopct="%1.1f%%", shadow=True, colors=colors, startangle=90, wedgeprops=wp, explode=explode, label='Percentage')

```

<Axes: ylabel='Percentage wise distrubution of feedback'>



```
1
2 #Feedback = 0
3 data[data['feedback'] == 0]['rating'].value_counts()
```

rating	count
1	544
2	128

```
1 #Feedback = 1
2 data[data['feedback'] == 1]['rating'].value_counts()
```

rating	count
5	1433
4	460
3	198

If rating of a review is 1 or 2 then the feedback is 0 (negative) and if the rating is 3, 4 or 5 then the feedback is 1 (positive).

```
1 #Distinct values of 'variation' and its count
2
3 print(f"Variation value count: \n{data['variation'].value_counts()}")
```

Variation value count:


variation	count
Tata-Tea-Gold-500g	58
Cinthol-Original-Soap-100g-Pack	40
Dettol-Liquid-Refill-Original-1500	40
Himalaya-Moisturizing-Aloe-Vera-200ml	40
Society-Tea-Masala-Jar-250g	40
..	..
Patanjali-UHT-Milk-1000-ml	10
Indiana-Frutti-Cherries-Frooti-Multicolor	6
Amul-Cow-Ghee-500ml	4
Tata-Tea-Premium-1-5kg	2
Patanjali-Ayurved-Ltd-CORO-NIL-Tablet	2

Name: count, Length: 122, dtype: int64

```

1 #Bar graph to visualize the total counts of each variation
2
3 data['variation'].value_counts().plot.bar(color = 'orange')
4 plt.title('Variation distribution count')
5 plt.xlabel('Variation')
6 plt.ylabel('Count')
7 plt.show()

```



```

1 #Finding the percentage distribution of each variation - we'll divide the number of records for each variation by total number of rec
2
3 print(f"Variation value count - percentage distribution: \n{round(data['variation'].value_counts()/data.shape[0]*100,2)}")

```

```

Variation value count - percentage distribution:
variation
Tata-Tea-Gold-500g                2.10
Cinthol-Original-Soap-100g-Pack   1.45
Dettol-Liquid-Refill-Original-1500 1.45
Himalaya-Moisturizing-Aloe-Vera-200ml 1.45
Society-Tea-Masala-Jar-250g       1.45
...
Patanjali-UHT-Milk-1000-ml        0.36
Indiana-Frutti-Cherries-Frooti-Multicolor 0.22
Amul-Cow-Ghee-500ml              0.14
Tata-Tea-Premium-1-5kg           0.07
Patanjali-Ayurved-Ltd-CORO-NIL-Tablet 0.07
Name: count, Length: 122, dtype: float64

```

```

1 #Mean rating according to variation
2 data.groupby('variation')['rating'].mean()

```

122 rows x 1 columns




Analyzing 'verified\_reviews' column This column contains the textual review given by the user for a variation for the product.

```
1 data['length'].describe()
```

	length
count	2763.000000
mean	222.631560
std	238.461988
min	2.000000
25%	57.500000
50%	159.000000
75%	312.000000
max	1866.000000


```
1 sns.histplot(data['length'],color='blue').set(title='Distribution of length of review ')
```

```
↳ [Text(0.5, 1.0, 'Distribution of length of review ')]
```




```
1
2 #Length analysis when feedback is 0 (negative)
3
4 sns.histplot(data[data['feedback']==0]['length'],color='red').set(title='Distribution of length of review if feedback = 0')
```

[Text(0.5, 1.0, 'Distribution of length of review if feedback = 0')]

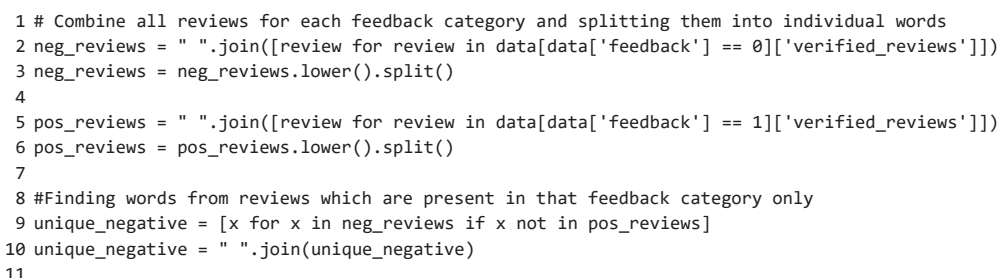
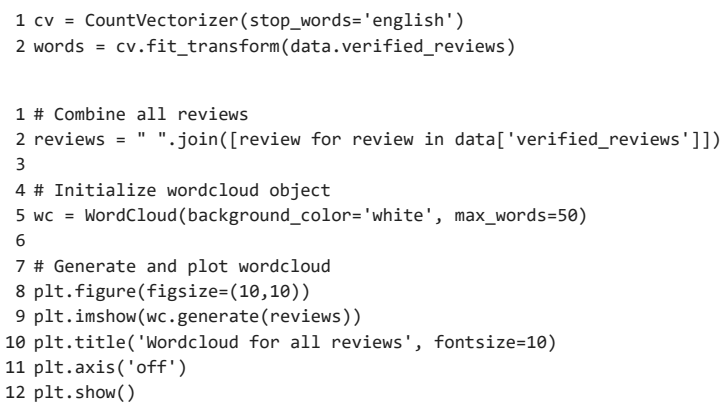


```
1 sns.histplot(data[data['feedback']==1]['length'],color='green').set(title='Distribution of length of review if feedback = 1')
```


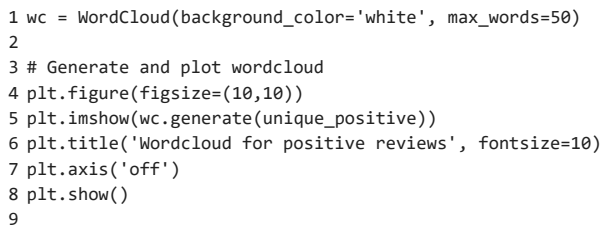
[Text(0.5, 1.0, 'Distribution of length of review if feedback = 1')]



```
1 #Lengthwise mean rating
2
3 data.groupby('length')['rating'].mean().plot.hist(color = 'blue', figsize=(7, 6), bins = 20)
4 plt.title(" Review length wise mean ratings")
5 plt.xlabel('ratings')
6 plt.ylabel('length')
7 plt.show()
```



```
1 wc = WordCloud(background_color='white', max_words=50)
2
3 # Generate and plot wordcloud
4 plt.figure(figsize=(10,10))
5 plt.imshow(wc.generate(unique_negative))
6 plt.title('Wordcloud for negative reviews', fontsize=10)
7 plt.axis('off')
8 plt.show()
```



1. Replace any non alphabet characters with a space
2. Covert to lower case and split into words
3. Iterate over the individual words and if it is not a stopword then add the stemmed form of the word to the corpus

<https://colab.research.google.com/drive/1-PS824UVoimhaly5an0MvWoAolQry958#scrollTo=sCYMvyCkhVXU&printMode=true>

```

4 review = re.sub('[^a-zA-Z]', ' ', data.iloc[i]['verified_reviews'])
5 review = review.lower().split()
6 review = [stemmer.stem(word) for word in review if not word in STOPWORDS]
7 review = ' '.join(review)
8 corpus.append(review)

```

```

1 #using Count Vectorizer to create bag of words
2
3 cv = CountVectorizer(max_features = 2500)
4
5 #Storing independent and dependent variables in X and y
6 X = cv.fit_transform(corpus).toarray()
7 y = data['feedback'].values

```

```

1 #Saving the Count Vectorizer
2 pickle.dump(cv, open("/content/sample_data/countVectorizer.pkl", 'wb'))

```

```

1 #Checking the shape of X and y
2
3 print(f"X shape: {X.shape}")
4 print(f"y shape: {y.shape}")

```

↗ X shape: (2763, 2500)  
y shape: (2763,)

```

1 #Splitting data into train and test set with 30% data with testing.
2
3 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3, random_state = 15)
4
5 print(f"X train: {X_train.shape}")
6 print(f"y train: {y_train.shape}")
7 print(f"X test: {X_test.shape}")
8 print(f"y test: {y_test.shape}")

```

↗ X train: (1934, 2500)  
y train: (1934,)  
X test: (829, 2500)  
y test: (829,)

```

1 print(f"X train max value: {X_train.max()}")
2 print(f"X test max value: {X_test.max()}")

```

↗ X train max value: 19  
X test max value: 16

```

1 #We'll scale X_train and X_test so that all values are between 0 and 1.
2
3 scaler = MinMaxScaler()
4
5 X_train_scl = scaler.fit_transform(X_train)
6 X_test_scl = scaler.transform(X_test)

```

```

1 #Saving the scaler model
2 pickle.dump(scaler, open("/content/sample_data/scaler.pkl", 'wb'))

```

```

1 #random Forest
2 #Fitting scaled X_train and y_train on Random Forest Classifier
3 model_rf = RandomForestClassifier()
4 model_rf.fit(X_train_scl, y_train)

```

↗ **RandomForestClassifier** ⓘ ?  
RandomForestClassifier()

```

1 #Accuracy of the model on training and testing data
2
3 print("Training Accuracy :", model_rf.score(X_train_scl, y_train))
4 print("Testing Accuracy :", model_rf.score(X_test_scl, y_test))

```

↗ Training Accuracy : 0.9963805584281282  
Testing Accuracy : 0.9650180940892642

```

1 #Predicting on the test set
2 y_preds = model_rf.predict(X_test_scl)

```

```


1 #Confusion Matrix
2 cm = confusion_matrix(y_test, y_preds)
3
4

```

```

1 cm_display = ConfusionMatrixDisplay(confusion_matrix=cm,display_labels=model_rf.classes_)
2 cm_display.plot()
3 plt.show()

```



```

1 #K fold cross-validation
2
3 accuracies = cross_val_score(estimator = model_rf, X = X_train_scl, y = y_train, cv = 10)
4
5 print("Accuracy :", accuracies.mean())
6 print("Standard Variance :", accuracies.std())

```

```


Accuracy : 0.956041343945302
Standard Variance : 0.015563172458390739

```

```

1 #Applying grid search to get the optimal parameters on random forest
2
3 params = {
4     'bootstrap': [True],
5     'max_depth': [80, 100],
6     'min_samples_split': [8, 12],
7     'n_estimators': [100, 300]
8 }
9
10
11 cv_object = StratifiedKFold(n_splits = 2)
12
13 grid_search = GridSearchCV(estimator = model_rf, param_grid = params, cv = cv_object, verbose = 0, return_train_score = True)
14 grid_search.fit(X_train_scl, y_train.ravel())

```



```

1 #Getting the best parameters from the grid search
2
3
4 print("Best Parameter Combination : {}".format(grid_search.best_params_))
5
6
7
8 print("Cross validation mean accuracy on train set : {}".format(grid_search.cv_results_['mean_train_score'].mean()*100))
9 print("Cross validation mean accuracy on test set : {}".format(grid_search.cv_results_['mean_test_score'].mean()*100))
10 print("Accuracy score for test set :", accuracy_score(y_test, y_preds))

```

```

Best Parameter Combination : {'bootstrap': True, 'max_depth': 100, 'min_samples_split': 8, 'n_estimators': 300}
Cross validation mean accuracy on train set : 97.66028955532575
Cross validation mean accuracy on test set : 89.80093071354705
Accuracy score for test set : 0.9650180940892642

```

```

1 #XgBoost
2 model_xgb = XGBClassifier()
3 model_xgb.fit(X_train_scl, y_train)

```

XGBClassifier

```

XGBClassifier(base_score=None, booster=None, callbacks=None,
               colsample_bylevel=None, colsample_bynode=None,
               colsample_bytree=None, device=None, early_stopping_rounds=None,
               enable_categorical=False, eval_metric=None, feature_types=None,
               gamma=None, grow_policy=None, importance_type=None,
               interaction_constraints=None, learning_rate=None, max_bin=None,
               max_cat_threshold=None, max_cat_to_onehot=None,
               max_delta_step=None, max_depth=None, max_leaves=None,
               min_child_weight=None, missing=None, monotone_constraints=None,
               multi_strategy=None, n_estimators=None, n_jobs=None,
               num_parallel_tree=None, random_state=None, ...)

```

```

1 #Accuracy of the model on training and testing data
2
3 print("Training Accuracy :", model_xgb.score(X_train_scl, y_train))
4 print("Testing Accuracy :", model_xgb.score(X_test_scl, y_test))

```

Training Accuracy : 0.9751809720785936  
Testing Accuracy : 0.9396863691194209

```

1 y_preds = model_xgb.predict(X_test)
2

```

```


1 #Confusion Matrix
2 cm = confusion_matrix(y_test, y_preds)
3 print(cm)
4
5
6
7 cm_display = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=model_xgb.classes_)
8 cm_display.plot()
9 plt.show()

```

```

[[151  32]
 [ 20 626]]

```



Decision Tree Classifier